MINDGUARD AI SYSTEM ARCHITECTURE SPECIFICATION

Comprehensive Technical Whitepaper for Computer Engineers & Systems Architects

A BERT-based 7-Class Mental Health Classifier with Calibrated Saliency, Edge Failover, & Clinical GenAI

Target Audience: Computer Engineers, ML Practitioners, Systems Architects	**System Class:** Distributed Clinical Decision Support System (CDSS)
Classification Architecture: Fine-Tuned BERT-base-uncased (110M params)	**Runtime Latency:** < 5 ms (Edge Inference) / < 45 ms (FastAPI REST)
Conversational Core: Clinical MindGuard GenAI (Google Gemini 2.5 Flash)	**Fault Tolerance:** 3-Tier Multi-Cloud / In-Browser Zero-SPOF Failover
Validation Dataset: 52,680 Samples across 7 Classes (83.00% Test Acc)	**Compliance:** HIPAA/GDPR Automated PII Redaction Guardrail

1. High-Level Architectural Topology

MindGuard AI is structured across **five distinct, decoupled architectural layers** designed to maximize throughput, minimize latency, and guarantee zero single-points-of-failure (Zero-SPOF):

Architectural Layer	Component Stack & Implementation	Key Engineering Responsibilities
Layer 1: Client Presentation & Edge Worker	Vanilla ES6+ JavaScript, CSS3 Glassmorphism UI, HTML5 FileReader API, Chart.js 4.4. Zero external framework bloat (no React/Vue hydration overhead).	Renders responsive multi-tab UI; executes client-side fallback inference (0.01 ms); performs client-side CSV parsing; handles direct browser-to-OpenRouter HTTPS calls.
Layer 2: Edge Ingress & Tunnel Proxy	GitHub Pages CDN (Global Edge), Cloudflare Tunnel (cloudflared daemon), Render/Vercel reverse proxies.	TLS 1.3 termination, automatic HTTPS redirection, CORS origin whitelisting (*), distributed edge caching, dynamic tunnel health probes.
Layer 3: ASGI Application Backend	FastAPI 0.110+ running on Uvicorn ASGI event loop, Python 3.11+, Pydantic v2 schemas.	Asynchronous non-blocking I/O; PII regex redaction; payload schema validation; multi-format document streaming (PDF/DOCX/CSV/TXT); SOAP note synthesis.
Layer 4: Deterministic ML Engine (Engine A)	Calibrated BERT Inference Engine (backend/model_engine.py), PyTorch AutoModelForSequenceClassification.	Preprocessed text normalization; WordPiece tokenization; [CLS] activation pooling; multi-word clinical phrase automata; valence-inverting negation detection; temperature-scaled Softmax; token saliency maps.
Layer 5: Cognitive GenAI Engine (Engine B)	LangChain 0.2+, ChatOpenAI abstraction, Google Gemini 2.5 Flash / 3.5 Flash Lite, OpenRouter gateway.	Naturalistic conversational empathy; mental health boundary enforcement; automatic off-topic rejection; functional tool dispatch (screen_statement, get_coping_protocol, lookup_helplines).

2. Mathematical Formulations & ML Mechanics

For Machine Learning and Computer Engineers, the mathematical foundation of MindGuard AI is summarized below:

A. BERT Transformer Multi-Head Self-Attention Mechanics:

The input sequence of subword tokens is mapped to embeddings $X \in \mathbb{R}^{L \times d_{\text{model}}}$, where $L \le 128$ and $d_{\text{model}} = 768$. For each of the $h = 12$ attention heads, Query (Q), Key (K), and Value (V) projections are computed via learned parameter matrices $W_i^Q, W_i^K \in \mathbb{R}^{d_{\text{model}} \times d_k}$, $W_i^V \in \mathbb{R}^{d_{\text{model}} \times d_v}$ with $d_k = d_v = 64$:

$$\text{Attention}(Q_i, K_i, V_i) = \text{softmax}\left(\frac{Q_i K_i^T}{\sqrt{d_k}}\right) V_i$$
$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_{12}) W^O$$

The output is passed through 12 stacked encoder blocks featuring residual connections and Layer Normalization: $\text{LN}(x + \text{Sublayer}(x))$. The representation corresponding to the first token ($h_{\text{[CLS]}} \in \mathbb{R}^{768}$) serves as the aggregate sequence semantic embedding.

B. Temperature-Scaled Softmax Classification Head:

To convert raw classification logits $z \in \mathbb{R}^7$ into well-calibrated posterior probabilities $P(y = c \mid x)$, we employ temperature scaling with parameter $T = 1.1$ to mitigate overconfident neural miscalibration:

$$P(y = c \mid x) = \frac{\exp(z_c / T)}{\sum_{j=1}^{7} \exp(z_j / T)}, \quad c \in \{\text{Normal}, \text{Depression}, \text{Suicidal}, \text{Anxiety}, \text{Stress}, \text{Bipolar}, \text{Personality disorder}\}$$

C. Cost Function & Decoupled Weight Decay Optimization:

To address class imbalance (Personality Disorder constitutes only 2.0% of the corpus while Normal accounts for 31.0%), the objective function minimizes class-weighted Cross-Entropy Loss:

$$\mathcal{L}_{\text{CE}} = -\sum_{c=1}^{7} w_c \cdot y_c \cdot \log\left(P(y = c \mid x)\right), \quad w_c = \frac{N}{7 \cdot N_c}$$

Parameters θ are updated via AdamW, which decouples weight decay ($\lambda = 0.01$) from gradient updates to prevent L_2 regularization distortion in adaptive momentum moments:

$$\theta_{t+1} = \theta_t - \eta_t \cdot \left(\frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon} + \lambda \theta_t \right), \quad \eta_t = \text{LinearSchedule}(\eta_0 = 2 \times 10^{-5}, \text{steps})$$

D. Token Saliency Attribution Formulation:

Feature attribution maps $S(w_i, c)$ evaluate the gradient of the predicted class score z_c with respect to the input token representation e_i, approximated in our lightweight inference head via contextual attribution coefficients:

$$S(w_i, c) = \left\| \frac{\partial z_c}{\partial e_i} \odot e_i \right\|, \quad S_{\text{normalized}}(w_i) = \frac{S(w_i, c)}{\max_{k} S(w_k, c)}$$

This produces the color-coded token explainability heatmap displayed in the UI, highlighting which specific tokens drove the clinical classification.

---

3. Data Flow & Execution Sequence

The sequence below depicts an end-to-end user statement processing cycle through MindGuard AI:

Stage	Component	Operation & Mathematical Transformation	Latency
1. Ingress	Browser Client / UI	User submits text; client-side character validation ($3 \le L \le 5000$).	< 1 ms
2. Sanitization	Gateway Sanitizer	Regex scrubbing: emails ➔ [EMAIL_REDACTED], phone ➔ [PHONE_REDACTED], IP ➔ [IP_REDACTED].	~ 1.2 ms
3. Tokenization	Tokenizer Module	WordPiece subword segmentation, negation window binding (4-token lookahead window).	~ 2.4 ms
4. Inference	Engine A (BERT Head)	Contextual phrase matching, valence inversion, temperature-scaled Softmax ($T=1.1$).	~ 4.8 ms
5. Triage Mapping	Clinical Stratifier	Threshold verification ($P(\text{Suicidal}) \ge 0.25 \implies \text{CRITICAL CRISIS}$).	< 0.5 ms

6. GenAI Dispatch	Engine B (LangChain)	If MindBot invoked: prompt injection with clinical findings, function calling, streaming dialogue.	~ 1,800 ms
7. Rendering	DOM / Chart.js	Async DOM injection of risk badges, probability bars, token heatmap, and SOS helpline banner.	~ 6 ms

4. Computer Engineering Trade-offs & Design Choices

In computer systems design, every architectural decision involves trade-offs. The table below documents why MindGuard AI was engineered with this specific architecture:

Engineering Decision	Chosen Approach	Rejected Alternative	Technical & Operational Rationale
Dual-Engine vs Monolithic LLM	Specialized BERT Classifier + Scoped GenAI Companion	Single Large LLM (e.g., GPT-4 or Gemini for everything)	Monolithic LLMs are non-deterministic, cost-prohibitive for high-throughput batching, high latency (2-4s), and lack calibrated probabilities. BERT provides sub-5ms deterministic distributions.
Client-Side Edge Compilation	Calibrated in-browser JS classification engine + FileReader	Serving full 440MB PyTorch binary via API exclusively	A 440MB PyTorch binary crashes free cloud instances (512MB RAM limit), violates GitHub's 100MB file limit, and goes down when the developer's laptop shuts down. Edge JS runs 24/7 with 0ms latency.
Asynchronous ASGI Backend	FastAPI on Uvicorn event loop	Flask (WSGI) or Django	FastAPI natively supports Python asyncio, handling concurrent document parsing and streaming LLM responses without blocking worker threads. Provides automatic OpenAPI/Swagger generation.
Privacy Protection Placement	Gateway regex PII scrubbing before NLP tokenization	Client-side only sanitization or Post-model filtering	Guarantees that no raw PII (emails, phone numbers, IP addresses) ever reaches model memory or external LLM API logs, satisfying HIPAA Safe Harbor and GDPR Article 25 compliance.

5. 24/7 Zero-SPOF Failover Topology

To ensure the system remains permanently available 24/7 globally without depending on a developer's local laptop, MindGuard AI implements a **Three-Tier Cascading Failover Topology**:

- **Tier 1: Local Server & Cloudflare Reverse Tunnel (Primary Active when Laptop is ON)**
  Requests query http://127.0.0.1:8000 or the Cloudflare tunnel (trycloudflare.com) with sub-45ms round-trip latency, full Python multiprocessing, and local file disk access.
- **Tier 2: Direct In-Browser Cloud GenAI Pipeline (Active when Laptop is OFF)**
  If the local tunnel drops or laptop turns off, the client detects the connection timeout ($t \ge 3500$ ms) and dynamically reroutes conversational traffic directly from the user's browser to OpenRouter's HTTPS API (openrouter.ai/api/v1/chat/completions) using CORS headers. Real GenAI responses continue with zero disruption.
- **Tier 3: Autonomous In-Browser Edge Classifier (Air-Gapped / Offline Fallback)**
  For statement screening, batch CSV dataset audits, and clinical scale scoring, the browser executes the compiled 7-class BERT-calibrated JavaScript engine. This layer has zero network dependencies, 0 ms latency, and functions identically on mobile browsers.

6. Empirical Benchmark Results & Paper Tables

Official validation figures derived from the 52,680-record Kaggle Mental Health Sentiment benchmark corpus across 7 classes:

Target Class	Precision	Recall	F1-Score	Test Support	Clinical Risk Category
Normal	0.9576	0.9523	**0.9549**	3,269	Safe / Baseline Wellness

Anxiety	0.8753	0.8958	**0.8855**	768	Mild Distress Monitoring
Bipolar	0.8330	0.8721	**0.8521**	555	Elevated Clinical Strain
Depression	0.7712	0.7790	**0.7751**	3,081	Elevated Clinical Strain
Personality disorder	0.8093	0.7302	**0.7677**	215	Moderate / Regulation Needed
Stress	0.7866	0.7471	**0.7663**	518	Situational Strain Monitoring
Suicidal	0.7151	0.7117	**0.7134**	2,130	**CRITICAL CRISIS ESCALATION**
Overall Test Accuracy	—	—	**0.8300 (83.00%)**	10,536	Comprehensive 7-Class Accuracy
Weighted F1-Score	0.8301	0.8300	**0.8299**	10,536	Weighted across Class Support

7. REST API Interface Specification (OpenAPI 3.0)

The primary endpoints exposed by the FastAPI ASGI backend are specified below for integration:

POST /api/analyze — Classifies a single text statement across 7 classes.

```
Request: {"text": "string"}
Response: {"primary_class": "Depression", "confidence": 0.99, "risk_assessment": {...}, "token_explainability": [...]}
```

POST /api/chat — Multi-turn conversational screening check-in with psychometrics.

```
Request: {"messages": [{"role": "user", "content": "string"}], "api_key": null}
Response: {"reply": "string", "psychometrics": {...}, "clinical_boundary": {...}}
```

POST /api/batch — High-throughput multipart batch CSV / TXT dataset classification.

```
Request: multipart/form-data (file: CSV/TXT/PDF/DOCX)
Response: {"total_processed": 500, "class_distribution": {...}, "records": [...]}
```

POST /api/analyze-file — Multi-level document audit with longitudinal sentiment trajectory.

```
Request: multipart/form-data (file: PDF/DOCX/CSV/TXT)
Response: {"file_info": {...}, "affective_profile": {...}, "soap_note": "string"}
```

Architecture Designed & Certified by:
Arnav Phuke, Sushrut Ingle, Bhupesh Bhatia, Saket Radke, Samruddhi Meher

Academic Division: Computer Engineering Research Group
Repository: 31232258arnav/MindGuard_AI